

SETU

सेतु — *Sanskrit for bridge.*

The War Room.

Everything you say, everything you show, and
every answer you will need when they push back.

Seven minutes on the clock. Eight minutes of questions from people who buy and build enterprise software for a living. This document is the complete script, the persuasion architecture, the corrected business model, the demo screenplay, and forty-four prepared answers — built against your final deck and against what your repository actually contains.



Read Part 2 tonight. It is the only part that can lose you the finale on its own.

Companion to SETU_Finale_Capgemini_T4PF_2026.pptx · Prepared 23 August 2026

Contents

Ten parts. Parts 2, 4 and 8 are the ones you rehearse from. The rest you read once and absorb.

Part 1	The room you are walking into — who judges, and on what	3
Part 2	Fix these before you present — deck errors and unsupported claims	5
Part 3	The story — narrative spine, the curb-cut reframe, real-world proof	8
Part 4	The script — every word, every pause, slide by slide	10
Part 5	The psychological playbook — sixteen moves and why each one works	15
Part 6	The business case — cost, price, margin, go-to-market, moat	17
Part 7	The demo — live runbook and the full video screenplay	21
Part 8	Jury defence — forty-four questions and their answers	24
Part 9	Pre-flight — the checklist and who does what	30
Part 10	Sources	31

“The market built tools that read the page. Nobody built an engine that finishes the task.”

That is your one sentence. If the jury remembers nothing else, they should remember that — and the fact that a clinician told you that you were wrong, and you rebuilt.

The room you are walking into

PART 1

You are not pitching a hackathon panel. You are pitching people who sign enterprise contracts. Everything below is drawn from Capgemini's own published material, not from guesswork.

WHO IS IN THE ROOM

Tech4Positive Futures is now in its third year. Hundreds of teams take part worldwide, country winners advance, and **the global finale is judged by members of the Capgemini Group Executive Committee** — the tier directly below the CEO. Expect a mixed India-US-Europe panel. Expect at least one person who has personally shipped enterprise software and at least one who runs a P&L.

This changes your register. Not "we built a cool thing." Rather: *here is a problem with a size, here is a product that runs, here is what it costs, here is who pays, and here is the specific thing I want from you.*

THE FOUR OFFICIAL CRITERIA

Capgemini states that Tech4Positive Futures solutions are assessed on four things. Learn these four words. They are your scoring rubric.

Criterion	What SETU has today	Where you are exposed
Quality of partnerships	A clinical reviewer, Dr. Prakhar Jain, who changed the build — and who has asked for a hospital version.	This is your weakest criterion. Capgemini's wording is specifically about NGO partnerships. You currently have zero. See the red box below.
Scalability	Genuinely excellent. Ten tools run locally at zero marginal cost; only the AI layer scales with spend.	Nothing. Lead with this.
Feasibility	Four surfaces deployed and running, 53 endpoints, a 32-point smoke test, a store-ready extension.	Only the claims that outrun the code. See Part 2.2 — this is the real risk.
Potential impact	Millions of students; a measurable outcome model; a clinical rewrite that proves you listen.	You have a product but no outcome data yet. Say so before they ask.

THE PARTNERSHIP GAP — FIX THIS IN 48 HOURS.

"Quality of partnerships" is an explicit, published judging criterion and your deck names no NGO. Every hour you spend on slide polish is worth less than one email to a disability NGO asking for a letter of intent to pilot. Candidates in India, in rough order of fit: **Madras Dyslexia Association** (Chennai, three decades of work in exactly this condition set), **Maharashtra Dyslexia Association**, **Dyslexia Association of India**, **NCPEDP** (national policy body for disability employment), **Action for Autism**, **Sarthak Educational Trust**, **ADAPT**. You do not need a signed contract. You need one named organisation who has seen the product and will say on paper that they will test it. One sentence on slide 6 — "[NGO] has agreed to run our first pilot cohort" — moves you from zero to scoring on a full quarter of the rubric. Verify any name and get written permission before you say it on stage.

THE COMPASS THEY JUDGE BY

Capgemini's CEO, Aiman Ezzat, has framed the group's technology philosophy as: *"technologies can enable progress for everyone if, first and foremost, they are designed by and for humans"* — the test being whether technology is **useful, accessible, and ethical**.

Those three words are a gift. SETU can claim all three literally, and one of them uniquely:

- **Useful** — it does not describe the page, it completes the task.
- **Accessible** — it is the product category itself.
- **Ethical** — you demoted your own headline feature when a clinician said it did not work, and your crisis path never touches a model. Very few teams can say either.

THE NUMBER YOU CAN HAND THEM

Capgemini publishes a digital-inclusion target: **10 million beneficiaries in underserved communities by 2030**. They report 8.4 million reached since 2018, and 955,000 in 2025 alone. The programme's three pillars are digital and green skills, pathways to employment, and applying technology to societal challenges.

USE THIS IN YOUR ASK. Their published inclusion programme does not currently name neurodiversity. You are not asking for charity — you are offering a route to a target they have already committed to publicly, in a population they are not yet counting. That is the single most executive-fluent sentence available to you, and it belongs in your closing ask:

"Your digital inclusion programme has reached 8.4 million people and is aiming at ten million by 2030. Neurodivergent learners are not named in it yet. We would like to be how you reach them."

WHAT THIS JURY REWARDS THAT A STUDENT PANEL DOES NOT

They reward	What that means for your seven minutes
Numbers with a denominator	"₹25 per active user per month" beats "very low cost" every single time. Never say "scalable" without a number behind it.
Deterministic behaviour	Enterprise buyers fear AI that surprises them. Your phrase is: <i>the agent will not act on a control the page did not report, and it never submits.</i>
Evidence of changing your mind	The clinical rewrite is worth more to this jury than any feature. Give it thirty seconds of stage time.
A specific, small, closeable ask	"One pilot campus this semester" is closeable. "Support us" is not.
Knowing your own weaknesses	Volunteering a limitation before they find it is the fastest credibility purchase available in the room.

Fix these before you present

PART 2

This is the most valuable part of this document. Two categories of problem: numbers on your slides that contradict each other, and claims in your master playbook that your repository does not support. Either one, found by a judge, damages everything else you said.

2.1 NUMBERS ON THE DECK THAT DO NOT AGREE

I read your final deck against your own cited sources. Six issues, in order of how likely a judge is to catch them.

#	On the slide now	The problem	What to put instead
1	Slide 5 headline: " Break-even at four campuses " with the ARR chart at ₹0.5 / 5 / 24.8 crore	Those figures were computed at ₹99 per student. You have raised the price to ₹199. The chart and the headline are now both wrong — and understated.	Break-even at two campuses . Chart: ₹1.0 Cr / ₹10 Cr / ₹49.8 Cr . Your new price makes the story <i>better</i> ; the slide just has not caught up.
2	Slide 2: " 5,04,080 " labelled " 0.1% of enrolment "	5,04,080 of 50 million is 1.0%, not 0.1%. The AISHE 2023–24 figure for students with benchmark disabilities is 54,080 , which is 0.1%. An order-of-magnitude error sitting next to its own contradiction.	54,080 • 0.1% of enrolment . The smaller number is also the more devastating one. Do not lose it.
3	Slide 2: " 12,000,000 – Alone in India " above a subline citing 8% of Indian children	12 million of 50 million is 24%, not 8%. The two numbers on the same slide disagree by a factor of three, and the source line underneath cites the 8% study.	See the decision box below — two defensible options, both stronger than the current mismatch.
4	Slide 2 headline: " 780 million Dyslexic, and 1 in 4 suffer ADHD GLOBALLY "	Two faults. The research says one in four people <i>with ADHD</i> also have dyslexia — you have reversed it. And " suffer " is the wrong word in a disability-inclusion pitch judged partly on dignity; the community's own language is "have".	"780 million people have dyslexia. One in four people with ADHD have it too."
5	Slide 2 fourth panel: " AI Integration / None – Implemented till date "	Reads at a glance as though SETU has no AI. It is directly beside the SETU panel. A judge scanning the slide will misread it.	Restore the parallel structure: AI CHATBOTS — "Answer questions about it." Or, if you want the "nobody has shipped this" point: AI THAT ACTS — "Nobody has shipped it."
6	Slide 3 title says three surfaces ; the speaker note underneath still says four surfaces . Slide 5 says " Free forever • 299 a month ".	The note is a leftover. The pricing line is missing its ₹ symbol, which on a global jury's screen could read as dollars — an 88× error.	Fix the note to "three surfaces". Write ₹299 . Also decide whether it is 30,000 or 50,000 lines and be able to say how you counted.

DECISION: what number goes on slide 2?

Option A — bulletproof. 45.0 million students in Indian higher education (AISHE 2023-24) $\times$ 8% pooled specific-learning-disorder prevalence (Scaria et al., meta-analysis of 8,133 children) = **3.6 million**. Registered: 54,080. The gap is **3.5 million**. Every number traces to a citation. Nothing on this slide can be knocked over.

Option B — bigger, still defensible. Widen the category from "specific learning disorder" to "neurodivergent", which independent estimates put at 15–20% of the population. 45.0 million $\times$ 20% = **9 million** — one student in five. The gap is **nearly 9 million**. You must then change the subline from "8% of Indian children" to "15 to 20 per cent of people are neurodivergent", and keep the 8% study as the sub-segment proof underneath.

Recommendation: Option B, with the subline changed. "One in five" is more memorable than "one in twelve", it is genuinely supportable, and it matches the global framing in your own headline. What is *not* defensible is the current slide, which shows a 24% number under an 8% citation. Whichever you choose, all three numbers on that slide — headline, subline, source — must tell the same story.

THE ONE-SENTENCE RULE. Before you present, take each number on each slide and say out loud: "That number comes from ___." If you cannot finish the sentence in under five seconds, the number comes off the slide. There is no third option.

2.2 CLAIMS IN YOUR MASTER PLAYBOOK THAT THE CODE DOES NOT SUPPORT

I searched your `chrome-extension/`, `backend/`, `frontend/` and `mobile/` trees for each technical claim in the Grand Finale Master Playbook. Several are not there. On a panel that includes people who have shipped browser extensions, "show me the WASM module" is an ordinary question, and there is no recovery from being unable to answer it.

Playbook says	Repository actually contains	Say this instead
Rust compiled to WebAssembly; sub-15 ms reflow on a web worker	No WASM, no Rust, no web worker anywhere in <code>chrome-extension/</code> . Fourteen plain JavaScript content scripts.	"The reading tools run entirely in the browser's own engine on the client CPU. No server call, no cost per use, and they keep working with the network off." — same benefit, and true.
Supabase PostgreSQL + pgvector RAG; "vector-indexed memory trails"	<code>mongodbService.js</code> — MongoDB, and it is optional. Without it everything falls back to local storage.	"MongoDB for persistence, and the product runs fully offline without it." The optionality is a <i>feature</i> — say so.
Next.js 15	React 18 + Vite (<code>frontend/vite.config.js</code>)	"React and Vite." Nobody will mark you down. Being caught renaming your stack is fatal.
Zero-Knowledge client-side AES-256 encryption of uploads	Not implemented. Identity is an anonymous per-device token in an <code>x-user-id</code> header. Journal entries are written to a local-only key and never mirrored to the database.	This is a stronger answer than the false one. "We do not encrypt-and-store. We mostly do not store. No name, no email, no account. The reflective journal never leaves the device, and the API keys never leave the server."
Blended cloud cost of \$0.015 per active user per month	At current model pricing the honest figure is roughly ₹25 (~\$0.28) — and today it is ₹0, because the model chain starts on free tiers.	Use ₹25 and ₹0. A judge who knows Gemini's per-token pricing can disprove \$0.015 in his head during your talk. ₹0 today is a better headline anyway.
Whisper API for speech	Sarvam (<code>speechService.js</code>) — Bulbul for synthesis, Saaras for transcription, eleven Indian languages.	"Sarvam, so it works in eleven Indian languages." This is a differentiator against every Western competitor. Do not trade it for a more famous name.
Expo SDK 57	Whatever is pinned in <code>mobile/package.json</code> . Check it tonight.	Say the version that is actually pinned, or just say "React Native with Expo".

WHY THIS MATTERS MORE THAN IT LOOKS.

A jury does not evaluate twenty claims independently. They sample two or three, and the result of that sample sets their confidence in the other seventeen. One unverifiable claim about WASM does not cost you the WASM point — it costs you the ₹25 point, the 53-endpoints point, and the clinical-review point, because now none of them have been checked. **Your real stack is impressive. It does not need help.**

2.3 THE HONESTY ASSET YOU ALREADY OWN

Your own source code contains something most teams would kill for. In `chrome-extension/content/eye-tracker.js`, the header comment reads:

“Honest description of the technique: this tracks *head position*, not pupil gaze... That is achievable without shipping a multi-megabyte CV model, and the UI says so rather than implying medical-grade eye tracking.”

You wrote a comment in your own codebase refusing to overclaim to *yourself*. If a judge presses you on whether you are exaggerating anything, that is your answer — and you can offer to show them the file. It is the most persuasive thirty seconds available to you in the entire Q&A, and it costs nothing because it is already true.

REHEARSE THIS LINE. *"We named a feature Gaze Scroll, and then wrote a comment in our own source code saying it tracks head position, not pupils, because we would not imply medical-grade eye tracking even to ourselves. I can show you the file. Everything else we have told you today was written under the same rule."*

The story

PART 3

Seven minutes is too short for a list and exactly right for a story. This part gives you the spine, the one reframe that converts this jury, and the real-world proof that makes it land.

THE SPINE: ONE STUDENT, ONE FORM, FIVE MINUTES

Do not present a population. Present a person, and return to her four times. The identifiable-victim effect is one of the most robust findings in decision research: people respond to one named individual far more strongly than to statistics about thousands, and the effect *reverses* when you lead with the number.

Use a specific student with a specific, time-bound, high-stakes task. Ours is:

Meera. Second year. Scholarship renewal portal. Last day to submit.

Not a generic "user". Not a 50-page portal in the abstract. A deadline, a consequence, and a name. Use the same name in the demo, so the jury watches *her* form get simplified, not *a* form.

Where	What Meera is doing
Slide 1	She opens the page. The letters move. Fourteen seconds later she closes the tab. (Your hook, made specific.)
Slide 2	She is one of nine million. And she has never been assessed, so officially she is one of nobody.
Slide 4	The demo <i>is</i> Meera. Say her name at the start of the demo and once more at the end.
Slide 6	She submits the form. On her own. That is the whole product, and the whole point.

THE REFRAME THAT WINS THIS ROOM: THE CURB-CUT EFFECT

Your biggest commercial risk in front of a P&L owner is the thought "*worthy, but niche*." There is one well-documented pattern that dissolves it, and it takes twenty seconds to deploy.

In the early 1970s, Berkeley, California cut the first ramps into its kerbs. They were built for a few hundred wheelchair users. Today the heaviest users are parents with prams, delivery workers with trolleys, travellers with suitcases, and anyone moving anything on wheels. Designers call it the **curb-cut effect**, and it is the most reliable pattern in the history of accessible design:

Built for	What it became	The lesson for your jury
Wheelchair users Kerb ramps, 1970s	Prams, luggage, trolleys, delivery	The accommodation became the default civic standard.
Deaf users Text messaging	The dominant communication mode on earth	A minority need became a majority behaviour.
Blind veterans Talking books, 1930s	The audiobook industry	An accommodation became a media category.
Hard-of-hearing viewers Closed captions	How most people watch video in public	The feature outgrew the population it was built for.
Arthritic hands OXO Good Grips	A mainstream premium kitchen brand	Designing for the hardest case produced the better product.

THE LINE: *"Every one of those started as an accommodation and ended as a market. Curb cuts were built for wheelchairs; today they are used by anyone with a suitcase. We are building the curb cut for attention."*

Deploy it as the bridge from the problem slide into the solution slide. It costs twenty seconds and it pre-answers the "isn't this niche?" question before anyone can ask it.

THE SECOND REFRAME: WHO IS ACTUALLY BROKEN

Your closing line — *"The interface was the disability. Not the person."* — is the strongest sentence in your deck. It is worth planting at the very start so the close pays it off. That is why the recommended edit to your hook ends on *"fixes the web, not the person."* Same idea, opening and closing. A jury that hears an argument returned to at the end scores it as complete rather than as a list.

WHY A CORPORATE JURY SHOULD CARE ABOUT NEURODIVERGENT TALENT

Keep this in your pocket for Q&A and for the employer half of your business model. Major employers already run structured neurodiversity hiring programmes — **JPMorgan Chase's Autism at Work now spans more than forty roles across nine countries**, and the *Harvard Business Review* case for neurodiversity as competitive advantage documents SAP, Hewlett Packard Enterprise and Microsoft reforming hiring to reach this talent and reporting gains in productivity, quality, innovation and engagement.

The commercial argument you can make to a services company in one sentence: *these firms have solved how to hire neurodivergent people, and none of them have solved how to make the internal tooling usable once they arrive*. That is the enterprise wedge, and Capgemini is exactly the kind of employer where it lands.

The script

PART 4

Written at roughly 145 words per minute — deliberately unhurried, because the subject is overload and your delivery is part of the argument. Read it aloud against a stopwatch four times before you change a single word.

THE CLOCK

Slide	Clock	Runs	Speaker	The one job this slide has
1	0:00–0:40	0:40	Speaker 1	Make them feel it. Name the product. Plant the reframe.
2	0:40–1:40	1:00	Speaker 1	Size the problem. Kill the three existing categories.
3	1:40–2:45	1:05	Speaker 1	Curb-cut pivot, then the architecture and the clinical rewrite.
4	2:45–5:15	2:30	Speaker 2	Prove it works. Five beats. Hands on the keyboard.
5	5:15–6:15	1:00	Speaker 2	Prove it survives outside a hackathon. Cost, price, break-even.
6	6:15–6:55	0:40	Speaker 1	The reframe returns. The ask. Then silence.

Finishes at 6:55 against a 7:00 hard stop. The one-minute warning at 6:00 lands mid-way through slide 5 — do not react to it, do not speed up, and above all do not skip the ask. If you are behind at 5:15, cut demo beat 5, not the business slide.

SLIDE 1 • 0:00 – 0:40 • SPEAKER 1 • THE HOOK

On screen: SETU. The chaos-to-clarity graphic on the right.

“Good morning respected judges and fellow innovators.

Imagine opening a website and trying to read an article... but the letters start moving, lines get mixed up, and your brain gets overwhelmed within seconds.

For millions of people with ADHD and dyslexia, this is not imagination — it is their daily reality.

Today we present SETU — **a solution that fixes the web, not the person.**

Setu is Sanskrit for *bridge*. It is a cognitive accessibility layer that sits between a page that was never designed for you, and the one thing you were actually trying to do.”

ACTION. Deliver the first three lines slowly and evenly — this is the only moment in seven minutes where you are allowed to be quiet and warm. Take a full second of silence after “*daily reality*”. Make eye contact with the most senior person in the room on “*fixes the web, not the person*”, and hold it for one beat. Do not smile through the hook; smile after it.

Your bracket, replaced — three options

1. "...a solution that fixes the web, not the person." **RECOMMENDED**

Seven words. Bookends your closing line exactly. Reframes the entire category before you have shown a single slide.

2. "...a solution that makes the web adapt to the mind, instead of the mind adapting to the web."

More explanatory, less punchy. Use if you want the mechanism stated up front.

3. "...a solution that turns the web from something they survive into something they can use."

The most emotional. Strongest if your delivery is naturally warm; weakest if it is naturally clipped.

SLIDE 2 • 0:40 – 1:40 • SPEAKER 1 • THE PROBLEM AND THE GAP

On screen: the two bars, the 95.9% panel, the four-category strip.

"There is a name for the moment I just described. Cognitive scientists call it the **Wall of Awful** — the instant executive function stops, because the interface is asking for more than working memory can hold. It is not a failure of intelligence. It is a failure of design.

Dyslexia affects around 780 million people worldwide. One in four people with ADHD have it too. Across the population, fifteen to twenty per cent of us are neurodivergent in some way.

In Indian higher education alone, that is about **nine million students**. One in five.

Now the number the system can actually see. **Fifty-four thousand and eighty**. That is every student in the country formally registered with a benchmark disability — all disabilities, not just these. Nought point one per cent.

[pause — two full seconds]

So nearly nine million students are getting through this alone. No assessment, no accommodation, and no idea that the problem was never them.

And the web is moving away from them, not towards them. WebAIM tested a million home pages this year: **95.9 per cent fail** accessibility standards. That is worse than last year — the first reversal in seven years.

Why has nothing solved it? Screen readers read the page aloud — the overwhelm is identical, just louder. Reading tools restyle the text — the task is exactly as hard. And AI chatbots answer questions *about* the page, then forget you the moment the tab closes.

The market built tools that read the page. Nobody built an engine that finishes the task."

ACTION. Speed up very slightly through the statistics — you are simulating the pile-up. Then stop dead for two seconds after "nought point one per cent" and let the invisible bar on the slide do the work; the silence is what makes them look at the graphic. On the last line, use your hands: left hand held low on "read the page", right hand raised and open on "finishes the task". The jury will remember the gesture longer than the sentence.

IF YOU ARE BEHIND. Cut the WebAIM sentences. Keep the two numbers, the pause, and the last line — they are the whole slide.

On screen: the engine diagram, three differentiators, the Dr. Jain quote band.

“Before I show you what we built — one pattern worth naming.

Kerb ramps were cut into pavements for wheelchair users. Today the people who use them most are parents with prams and travellers with suitcases. Text messaging was built for deaf users. Captions were built for the hard of hearing, and are now how most people watch video on a train. **Every one of those started as an accommodation and ended as a market.** We are building the curb cut for attention.

SETU is one engine behind three surfaces. A Chrome extension that reshapes the page you are on. A web workspace that turns a document into something you can actually hold in your head. And an Android app, so the task survives leaving your desk. Fifty-three endpoints, nine cognitive modes, eleven Indian languages — deployed and running right now.

Three things make it hard to copy.

One — it is local-first. Ten of the reading tools never touch a server. They cost us nothing per use, and they keep working with the network off.

Two — it degrades instead of failing. Every AI feature draws a usable answer off the page in under a second, then upgrades it when the model replies. A tool that shows a spinner and then an error is a tool nobody opens twice.

Three — and this is the one I would ask you to remember. On the seventeenth of August we took SETU to Dr. Prakhar Jain. He told us that bigger fonts and reading aids do not solve dyslexia, because the difficulty is comprehension, not eyesight. He was contradicting our headline feature.

So we demoted it. And we rebuilt: a Numbers mode for dyscalculia that uses countable objects inside a story, a picture mode that carries no text at all, and a crisis path that never reaches a model.

We did not add a clinician to a slide. We let one change the product.”

ACTION. The curb-cut section should sound like an aside — slightly slower, slightly conversational, as though you are letting them in on something. Then snap back to pace for the architecture. On “*we demoted it,*” slow right down. This is the single highest-scoring sentence in your seven minutes and most teams would never say it out loud. Speaker 2 walks to the laptop during the final line — do not narrate the handover.

IF YOU ARE BEHIND. Cut the curb-cut paragraph and redeploy it in Q&A when someone asks about market size. Never cut the clinical rewrite.

On screen: switch to the browser. The deck slide is a five-second anchor only.

“This is Meera. Second year, and it is the last day to renew her scholarship. Let us do it with her.”

BEAT 1 • The portal (0:00 – 0:30)

“Fourteen fields. Three warnings. No indication of what to do first. For most of us this is annoying. For Meera this is where the task ends — and the deadline is today.”

BEAT 2 • Alt + 3 (0:30 – 1:05)

“One keystroke. Three steps, in under a second — and notice what did *not* happen. Nothing left this browser. No server, no API key, no cost to us at all. This is one of the ten local tools, and it is why our margin story works.”

FALLBACK. If the chunker stalls, press **Alt+F** for Focus Mode instead — it strips the page live, locally, and makes the identical point. Say: *“Focus Mode, same idea — the page gets quieter before anything gets smarter.”*

BEAT 3 • Commander (1:05 – 1:45)

“Now the agent. I ask it to fill the form. It reads the page's own reported controls, plans the steps, fills them — and then it **stops**. It will not submit. Anything irreversible needs a human to say yes.

That is not a policy we wrote in a slide. It is asserted in our test suite: no step may target a control the page did not report, every submit step carries a confirmation requirement, and no personal data is ever invented to fill a field.”

ACTION. This is the trust beat, and it is the answer to the question every technical judge is already forming. Slow down. Let the cursor sit on the unsubmitted form for a full two seconds before you speak again. Silence over a stopped agent is more convincing than any sentence you could say over it.

BEAT 4 • The voice (1:45 – 2:15)

“Dr. Jain told us mind maps only help if they carry no text — or if the text is spoken the moment the cursor reaches it. So: hover a branch, and it speaks. In Tamil. Eleven Indian languages, and the data schema underneath stays in English, so nothing silently breaks when the language changes.”

FALLBACK. If the voice is slow, the browser's own engine takes over automatically. Say it out loud: *“— and there is the fallback. The user never hears silence.”* The failure demonstrates the feature. This is the best possible thing that can go wrong.

BEAT 5 • The handoff (2:15 – 2:30)

“Same student, same task, on the phone. It picks up mid-sentence. Working memory is the thing that is actually impaired here — so SETU carries the context, and Meera does not have to.”

CLOSE THE DEMO ON HER NAME. *“Meera submits the form. On her own. That is the entire product.”* Then return to the deck.

SLIDE 5 • 5:15 – 6:15 • SPEAKER 2 • COST, MODEL, SCALE

On screen: the ₹25 cost stack, the three pricing tiers, the growth chart, the Dr. Jain band.

“Now the part that decides whether this survives past this room.

One active user costs us **twenty-five rupees a month**. Eight for AI reasoning, fifteen for the natural voice, two for storage — and nothing at all for the ten local tools, because they never call us.

Today it costs us **zero**. The model chain starts on free tiers and ends in a deterministic offline engine, so our infrastructure spend to date is nil. Four production surfaces, built by five final-year students, for the cost of our time.

We sell to campuses at **₹199 per student per year, licensed campus-wide** — because the university cannot possibly know which students need it. Ninety-eight per cent of them were never assessed. Campus-wide is not a pricing trick; it is the only honest way to price an invisible population.

Year one costs us eighteen lakh to build and run. **Two campuses cover it**.

And this is not theoretical. Dr. Prakhar Jain — the specialist who reviewed us — has asked us to build a version of SETU for his hospitals. That came to us. Before we launched.”

ACTION. This is the most number-dense minute in the pitch, so slow your delivery rather than speeding it up. Land **₹25**, **₹0**, and **two campuses** as three separate, deliberate beats with a half-second of air around each. The one-minute warning bell falls somewhere in here — ignore it completely; your plan already accounts for it.

Do not present this slide until the chart and the headline have been corrected to your new ₹199 pricing. See Part 2.1, item 1. At ₹199 you break even at two campuses, not four, and year three is ₹49.8 crore, not ₹24.8 crore. Presenting stale numbers next to your new price is the one thing on this slide a judge is certain to catch.

SLIDE 6 • 6:15 – 6:55 • SPEAKER 1 • THE CLOSE AND THE ASK

On screen: the closing line, the three pilot metrics, the ask.

“Every one of those nine million students has been told, in one way or another, that they are the problem.

They are not. **The interface was the disability. Not the person.**

Your digital inclusion programme has reached 8.4 million people and is aiming at ten million by 2030. Neurodivergent learners are not named in it yet. We would like to be how you reach them.

Three asks. **One pilot campus**, through a Capgemini client, this semester. **One introduction** to a university disability support office. And **back SETU in this finale** — we will publish the numbers either way.

Thank you.”

ACTION. Land “*The interface was the disability. Not the person.*” and then stop for a full second before moving on. Say the three asks on your fingers — one, two, three — because a jury that can count your asks can grant one. Then say thank you, stand still, and **say nothing else**. Do not fill the silence. Do not thank the organisers again. Do not add “and we’re happy to take questions.” The last sentence they hear is the one that follows them into scoring.

The psychological playbook

PART 5

Sixteen moves. None of them requires overstating anything — your material is strong enough that the honest version is also the persuasive one, which is the only kind of persuasion that survives eight minutes of cross-examination.

THE ORGANISING MOVE

This is a pitch about cognitive overload, so the pitch itself is not permitted to cause any. One idea per slide, numbers at fifty point, almost no prose, and a colour palette taken directly from the product's own design system. A jury sitting through a dozen dense decks will feel the difference before they can articulate it — and if anyone asks, the honest answer is the best one:

“We built the deck the way we built the product.”

THE SIXTEEN MOVES

Move	What you actually do	Why it works
1. Sensory contrast	Speed up slightly through the slide 2 statistics. Then drop your pitch, slow down and lower your volume the moment Focus Mode engages in the demo.	You induce a mild version of the problem, then relieve it. The jury <i>experiences</i> the product rather than hearing about it. This is the strongest single technique available to you.
2. The identifiable victim	Name her. Meera, second year, scholarship deadline today. Return to her four times.	People respond far more strongly to one named person than to statistics — and the effect reverses if you lead with the number. One face, then the population.
3. Costly signalling	"A clinician told us our headline feature does not work. We demoted it."	Admitting error is expensive, which is exactly why it is believed. Nothing else in seven minutes buys this much credibility this cheaply.
4. The pratfall effect	Volunteer your one real weakness: "We have a working product and no outcome data yet. That is what the pilot is for."	Competence plus a disclosed flaw is rated higher than flawless competence. A team with no admitted weakness reads as a team that has not looked.
5. Specificity	54,080. 95.9%. ₹199. 48,246 colleges. Never "about" anything.	Precise numbers read as measured; round numbers read as invented. Precision is a claim about your process, not just your data.
6. Let them do the maths	The 54,080 bar is drawn to true scale, so it is almost invisible. Do not explain it. Pause instead.	A conclusion the audience reaches themselves is held far more firmly than one they are handed.
7. Loss framing	"95.9% fail — worse than last year. The first reversal in seven years."	Deterioration mobilises where a static bad number does not. It converts a problem into a deadline.
8. The two-second pause	After "nought point one per cent." After "we demoted it." After the closing line. And before answering any aggressive question.	Silence signals control. Rushing to fill it signals the opposite. In Q&A a two-second pause reads as consideration; an instant answer reads as a script.
9. Peak-end	The demo (highest sensory moment) sits in the middle. The reframe (highest emotional moment) sits at the very end.	People remember the peak and the ending, and almost nothing in between. Engineer both deliberately.
10. Bookending	Open on "fixes the web, not the person." Close on "the interface was the disability, not the person."	An argument returned to at the end is scored as complete. An argument that only accumulates is scored as a list.

Move	What you actually do	Why it works
11. Executive mirroring	Use their vocabulary: <i>deterministic, human-in-the-loop, marginal cost, gross margin, compliance exposure, institutional procurement</i> . Never "hack", "cool", "we just".	Register signals category. You want to be evaluated as a company, not as a student project, and the vocabulary decides which frame they use.
12. Their own KPI	"8.4 million reached, ten million by 2030, and neurodiversity is not named in it yet."	The most persuasive thing you can offer a large organisation is progress against a number they have already published. It converts your ask from a request into a contribution.
13. Category reframe	The curb-cut effect. Accommodations become markets.	Pre-empts "worthy but niche" before it forms. You cannot argue someone out of that thought once they have had it — you have to prevent it.
14. Authority, twice, differently	Dr. Jain appears once as the reason the product changed, and once as unsolicited demand. Never as a generic endorsement.	One name doing two different jobs reads as depth. The same name doing the same job twice reads as thin.
15. The three-item ask	Pilot campus, then an introduction, then a vote. Say them on your fingers.	Descending order of cost to them. Anyone who declines the first will feel the pull of the second. A countable ask is a grantable ask.
16. Terminal silence	After "thank you", stop. Stand still. Say nothing.	Trailing words after a strong close overwrite it. The most disciplined thing you will do all day is the two seconds after your last sentence.

BODY AND VOICE — THE PHYSICAL CHECKLIST

Do

- Feet planted, shoulder width. Move only between slides, never mid-sentence.
- Hands visible and open. Low hand for "reads the page", raised hand for "finishes the task".
- Rotate eye contact across three jurors, holding each for one complete sentence.
- Drop your pitch at the end of key sentences. Rising intonation turns a statement into a request for approval.
- Breathe out before you start speaking, not in.
- When the other speaker is talking, look at the screen or at them — never at your notes.

Do not

- Read the slides. Every number on screen exists so you do not have to say it twice.
- Apologise — not for latency, not for prototype status, not for being students. State constraints as design decisions, because that is what they are.
- Say "just", "basically", "kind of", "we tried to".
- Fill a pause. The pause is doing work.
- Debug live. If something breaks, say the fallback line and keep moving.
- Speed up when the six-minute bell rings. Your plan already accounts for it.

THE ONE RULE THAT OUTRANKS ALL SIXTEEN. Every technique above amplifies whatever is underneath it. Applied to a claim you cannot support, they amplify the damage instead. Run Part 2 first, then Part 5.

The business case

PART 6

Rebuilt against your new ₹199 price and against what the code actually costs to run. Every figure here traces to a source in Part 10 or to a line in your repository.

6.1 WHAT IT HAS COST TO BUILD

₹0

Infrastructure spend to date. Render, Vercel, MongoDB Atlas and OpenRouter free tiers, with a deterministic offline engine underneath for when they rate-limit.

4

Production surfaces — extension, web workspace, Android app, shared engine — built and deployed by five final-year students.

This is a genuinely strong statement and most teams cannot make it. It also sets up the margin argument: a product whose core costs nothing to run was *designed* that way, not lucky.

Year one, productionised

Line	Basis	Cost
Two engineers, full time	12 months × ₹60,000	₹14,40,000
AI, voice and infrastructure at ten campuses	~700 monthly-active × ₹25 × 12	₹2,10,000
Store fees, accessibility audit, trademark clearance, legal	One-off	₹1,50,000
Total year one		₹18,00,000

6.2 WHAT ONE USER COSTS TO SERVE

Line	Basis	Per active user / month
AI reasoning	Gemini 2.5 Flash at \$0.30 per million input and \$2.50 per million output tokens; roughly 30 operations at 3,000 in / 800 out	₹8
Natural voice	Sarvam Bulbul v2 at ₹15 per 10,000 characters; roughly 10,000 characters, with both-ends caching	₹15
Storage and sync	MongoDB, amortised	₹2
Ten local tools	Run in the browser. They never reach our servers.	₹0
Total	And today, ₹0 — the model chain begins on free tiers.	₹25 (≈ \$0.28)

Retire the \$0.015 figure from the master playbook. At published model pricing it is not reachable, and a judge who knows token economics can disprove it mentally while you are still speaking. ₹25 with a visible derivation is more persuasive than \$0.015 with none — and "today it is ₹0" is a better headline than either.

6.3 PRICING, AND WHY IT IS SHAPED THIS WAY

Tier	Price	The argument
Campus	₹199 / student / year campus-wide	The university cannot identify who needs it — India registers 0.1% of students as disabled against a far higher prevalence. A per-seat licence would only reach students already being supported, which is the group that needs you least. Campus-wide is the only model that reaches the rest, and it works commercially <i>because</i> most licensed students only ever use the free-to-serve local tools.
Employer	₹499 / employee / year	Paid from accessibility and DEI budgets that already exist. Compliance drivers: the European Accessibility Act, in force since June 2025, and the RPwD Act and GIGW in India. Major employers have solved <i>hiring</i> neurodivergent talent; none has solved making the internal tooling usable once they arrive.
Individual	Free forever ₹299 / month premium	The free tier genuinely costs nothing to serve, so it is a growth loop rather than a subsidy. Premium buys natural voice, AI features and cross-device sync.

If a judge asks why ₹299 a month retail against ₹199 a year institutional: "Campus licences are bought for five thousand seats of which about twelve per cent activate, so the effective price per *active* student is around ₹1,650 a year — about ₹138 a month. Against ₹299 retail, that is a 2.2× volume discount, which is ordinary enterprise pricing." Have this ready. The apparent 18× gap is the most obvious thing to attack on that slide.

6.4 UNIT ECONOMICS OF ONE CAMPUS

One 5,000-student campus, one year	Amount
Licence revenue — $5,000 \times ₹199$	₹9,95,000
Cost to serve — 12% monthly-active (600) $\times ₹25 \times 12$	– ₹1,80,000
Allocated infrastructure	– ₹40,000
Gross margin	₹7,75,000 · 78%
Break-even — $₹18,00,000 \text{ year-one cost} \div ₹9,95,000$	2 campuses

Note for your own confidence: 78% is a normal software gross margin and it is high *because* the local-first architecture keeps most users off your servers entirely. If a judge probes the margin, that is the answer — and the extra headroom is what funds the permanently free individual tier.

6.5 SCALE — CORRECTED FOR ₹199

Year	Institutions	Students licensed	ARR
Year 1	10	50,000	₹1.0 crore
Year 2	100	500,000	₹10.0 crore
Year 3	500	2,500,000	₹49.8 crore ($\approx \$5.7\text{M}$)

Five hundred institutions is roughly **one per cent** of India's 48,246 colleges. Framing the year-three number as one per cent of the market is far more credible than framing it as a large number in isolation — it invites the jury to check the arithmetic and find it conservative.

Market sizing, done honestly

	Figure	Basis
TAM	\$21.4B → \$48.8B by 2035 8.6% CAGR	Global assistive technology market. Cognitive and adaptive software is a smaller slice, but the same analysis flags it as strategically significant for education and workplace accessibility, with high switching costs and deep integration.
SAM	≈ ₹895 crore (~\$100M)	Indian higher education fully licensed: 45 million students × ₹199. This is your honest serviceable market, not a global fantasy.
SOM	₹49.8 crore in three years	500 institutions — 1% of colleges.

Drop the "\$13 trillion TAM" from the master playbook. That figure is the global *disability economy* — the household spending power of disabled people and their families. It is not an addressable software market, and to a Group Executive Committee it will read as either careless or inflated. A \$100 million Indian SAM that you can defend line by line is worth more than a \$13 trillion number that you cannot.

6.6 THE SEVEN UNIQUE SELLING IDEAS, RANKED FOR THIS JURY

#	The idea	Why it scores here
1	It finishes the task	Every competitor optimises reading. You are the only one that acts on the page. This is category creation, not feature competition — and category creators get remembered.
2	Local-first, zero marginal cost	Ten tools never touch a server. This is the entire margin story and it is architecturally true, not aspirational. Scores directly on "scalability".
3	Degrades, never fails	Free tier → paid → deterministic offline engine, with something usable rendered off the page in under a second. Enterprise buyers care more about predictable failure than peak capability.
4	Rewritten on clinical evidence	You demoted your own flagship feature. Scores on impact, on ethics, and on the jury's private question — <i>will these people listen to a customer?</i>
5	Eleven Indian languages	No Western competitor will build this. It is a durable regional moat and it maps onto Capgemini's India footprint exactly.
6	Human-in-the-loop by construction	The agent never submits, and that constraint is asserted in the test suite rather than promised in a slide. This is the answer to the AI-safety question before it is asked.
7	A crisis path that never reaches a model	Detected server-side, answered from a fixed reviewed script with real helplines, never sampled, never scored. Almost nobody builds this. It signals that you thought about the worst case, not just the demo case.

6.7 THE MOAT

Any single feature here is copyable. The combination compounds, and that is the argument:

- **Behavioural detection** tells you when someone is struggling — which makes the intervention timely rather than annoying.
- **A persistent cross-surface profile** makes each intervention better than the last, and it moves with the person between browser, desktop and phone.
- **Structured task breakdown** is the part that produces the outcome the institution is buying.
- **Eleven Indian languages** makes the whole thing usable in the market where you will win first.

Competitors have one or two of these. None has all four, and the loop tightens with use: the more consistently SETU helps someone act, the more accurately it can reduce friction next time.

The second moat is distribution, and it is the one that actually holds. A disability support office that has run a semester pilot, has your outcome numbers, and has trained its staff on your product does not re-procure casually. Institutional accessibility software has notoriously high switching costs. Your defensibility ends up being the relationship, not the repository.

6.8 RISKS, AND HOW YOU ANSWER THEM

Risk	Your answer
A large platform ships this	Immersive Reader has existed for years and 95.9% of the web still fails. They ship reading, not task completion, and they will not build eleven Indian languages with a dyscalculia mode for a market this size.
Free model tiers disappear	Already priced. ₹25 per active user per month <i>is</i> that scenario, and the fallback engine below it is deterministic and free.
Voice cost rises	Clips cache on both ends, the browser voice is a working fallback at zero cost, and TTS is the most competitive segment in AI right now — the price risk points down.
Activation is higher than 12%	Then margin compresses and price rises. At 20% activation the campus price moves to about ₹299. The model still works; we would rather discover that on one campus than five hundred.
Long institutional sales cycles	We start at the disability support office, not central procurement. They have a statutory mandate, a budget line and a waiting list — a semester pilot is a far shorter decision than an IT purchase.
The name "SETU"	There are existing SETU marks in India in Class 9 and Class 42. We would run a clearance search before commercial launch and we are not attached to the word. The product is the bridge.
The team graduates	Honest answer: a signed pilot changes the decision. That is precisely why the ask is a pilot rather than a prize.

6.9 THE PARTNERSHIP PLAY

Partnerships are a quarter of the published rubric and currently your weakest quarter. Three moves, in order of how fast they can be done:

- 1. An NGO letter of intent, before the finale.** One named disability organisation who has seen the product and will state on paper that they will test it with a cohort. This is a two-day task and it is worth more than any slide edit.
- 2. A university disability support office.** Named, with a target cohort size and a semester. "A 250-student cohort at [institution] this semester" is concrete; "we will pilot with universities" is not.
- 3. Capgemini itself, framed as co-innovation rather than charity.** Their employee resource groups as a proving ground, and their published inclusion target as the destination. This is the ask that makes the jury a participant rather than a judge — and people evaluate favourably things they imagine themselves inside.

The demo

PART 7

Two and a half minutes of your seven. This part covers the live runbook, the failure table, and a complete shot-by-shot screenplay for the recorded version.

7.1 LIVE OR RECORDED?

Record the video regardless. Then decide on the day. Live scores higher when it works, because a jury that has sat through nine polished videos is starving for something real. Recorded scores higher when the venue's network is unknown and your engine sleeps on a free tier.

The hybrid, and our recommendation: run the video for beats 1, 2 and 4 (visually rich, timing-critical, unforgiving if slow), then go live on your own laptop for beat 3 — the agent filling the form and stopping before submit. Twenty seconds of live, at the exact moment the jury most wants to believe you. Say it out loud: *"That was recorded. This next part is not — let me do it here."* The contrast is the point.

7.2 PRE-FLIGHT

The night before

- Render the video and put it in the last tab, muted, full screen ready.
- Run `npm run smoke` in `backend/`. Thirty-two checks. Know the result before the room does.
- Pin exactly three tabs: portal, Sanctuary, video. Close everything else.
- Browser zoom 125%, display 1920×1080. Jurors eight metres back cannot read your default size.
- Phone charged, mirroring tested on the venue display, same task already open.

Five minutes before

- Wake the engine — the free tier sleeps and takes up to a minute. Hit the health endpoint twice.
- Confirm the popup status pill is **green**. Amber means no AI key; pink means unreachable.
- Press **Alt+X** so every feature is off and beat 2 starts clean.
- Mobile hotspot on and tested as a fallback to venue Wi-Fi.
- Notifications off. Sleep and screensaver disabled. Power connected.

7.3 IF SOMETHING BREAKS

Failure	Say this, without pausing	Do this
Engine unreachable	"That is the free tier asleep — which is exactly why the reading tools do not depend on it."	Press Alt+F. Focus Mode is fully local.
AI chunking slow	"It has already drawn the steps off the page. The model is only improving them."	Keep talking. The local render is already on screen.
Voice fails	"And there is the fallback — the user never hears silence."	Nothing. The browser voice takes over by itself.
Agent misfires	"It will not act on a control the page did not report. That is the safety rule doing its job."	Move to beat 4. Never retry live.
Phone will not mirror	"Take my word for the handoff, or catch me afterwards and I will show you."	Skip beat 5. Go to slide 5 early.
Total network loss	"Everything you are about to see runs offline anyway."	Beats 1, 2 and 5 still work.

Notice that four of those six failures are answered by an architectural decision you already made. That is not luck — it is what local-first with a deterministic fallback buys you. If one of them happens, say so explicitly. A failure that demonstrates your design is better than a demo that goes perfectly.

7.4 THE VIDEO — FULL SCREENPLAY

Two minutes thirty. No music, or something very quiet — the subject is sensory overload and a driving soundtrack contradicts your entire thesis. One voice, calm, the same person who narrates live. Screen recording at 1920×1080, 30fps, cursor highlighting on.

Time	On screen	Callout badge	Voice-over
0:00–0:12	Cold open. A cursor hovering, motionless, over a dense scholarship portal. No movement for three full seconds. Then the tab closes.	— none —	"This takes fourteen seconds. Watch." Then silence until 0:12. The stillness is the shot.
0:12–0:35	Reopen the same portal. Slow pan across fourteen fields, three red validation warnings, a countdown banner.	THE STATUS QUO	"This is Meera's scholarship renewal. Fourteen fields, three warnings, and today is the last day. For a reader with dyslexia the page is not difficult. It is closed."
0:35–0:58	Extension icon click. Clutter dissolves. Bio-Reading fixation weighting appears on the text; the line ruler tracks down the page.	RUNS IN THE BROWSER · ₹0 PER USE	"One click. The page gets quieter. Bio-Reading anchors each word's fixation point and the ruler holds the line. All of this runs on her own machine — nothing has been sent anywhere, and it costs us nothing."
0:58–1:20	Cursor begins moving erratically — reversals, rapid scroll thrash, three clicks on inert text. Screen dims gently into a four-second box-breathing animation.	OVERWHELM DETECTED · LOCALLY	"Now watch the cursor. Reversals, scroll thrash, clicks on things that are not buttons. SETU reads that pattern — in the browser, from movement alone, no camera, nothing leaving the device — and does not shout. It offers a breath." Drop your volume here. Slow down. Let four seconds of the animation play unnarrated.
1:20–1:40	Card: "Shall I break this into three steps?" Click yes. The dense form resolves into three linear action cards.	THREE STEPS · UNDER ONE SECOND	"Three steps. Under a second, because they were built from the page itself before any model was called."
1:40–2:05	Microphone tap. Voice command in the user's own language. Fields populate one by one. The cursor arrives at Submit — and stops. A confirmation prompt appears. Hold on the unpressed button for two seconds.	IT FILLS. IT DOES NOT SUBMIT.	"She asks it to fill the form. It reads the page's own reported controls and fills them — and then it stops. It will not submit. Anything irreversible needs a human." Say nothing over the final two seconds. Let the unpressed button sit there.
2:05–2:22	Split screen. Phone appears beside the desktop. Same checklist, same position, mind map branch speaking aloud in Tamil with the caption visible.	ELEVEN INDIAN LANGUAGES	"She leaves her desk. The task comes with her — same step, same place, and the map speaks in her language."
2:22–2:30	Back to the portal. Submit is pressed. A confirmation screen. Hold. Fade to the SETU wordmark.	— none —	"Submitted. On her own. That is the entire product."

Capture notes

- **Use a real portal, not a mock.** Judges can tell, and a real government or university form is a stronger claim than any staged page. Blur any genuine personal data in post.

- **Do not speed up the footage.** If a step is slow, cut to the result — never ramp the playback. Sped-up screen recording is the most common tell that a demo is not real.
- **Record the overwhelm cursor deliberately.** Rehearse the erratic movement until the detector fires reliably, then keep that take.
- **Caption everything.** An uncaptioned accessibility video will be noticed by this jury specifically, and it is an unforced error.
- **End on the submitted confirmation, not on your logo.** Let the wordmark come after. The last frame should be the outcome, not the brand.

Jury defence

PART 8

Forty-four questions across eight categories. Eight minutes of Q&A is longer than your presentation and it is where the finale is actually decided. Answer in three sentences and then stop — a short confident answer scores higher than a complete one.

A • IMPACT AND EVIDENCE

A1. What evidence do you have that this actually helps?

Two kinds. Structured task breakdown and single-step presentation are established practice in executive-function coaching and special education — we did not invent the intervention, we automated it. And we ran the product past a clinician on the seventeenth of August, which changed the build substantially. What we do not yet have is outcome data, and that is exactly what the pilot is for.

Established intervention, automated. Outcome data is what the pilot buys.

A2. What did the clinical review actually change?

More than we expected. Dyslexia was separated into three conditions in onboarding, with a new Numbers mode for dyscalculia that renders countable objects inside a story. Bio-Reading was demoted from headline feature to an honestly labelled comfort option. Mind map branches now speak on hover and on keyboard focus, and a picture mode strips the prose entirely. A reward system replaced text reports for ADHD, a parking lot was added for impaired working memory, and a reflective listener was added because roughly forty per cent of this population also experiences anxiety or depression.

He contradicted our flagship. We rebuilt around him.

A3. How will you measure whether it worked?

Three numbers, all measurable inside one semester: time to first action, task completion rate on dense multi-step forms, and the proportion of tasks completed without asking another person for help. The third one is the real outcome — independence, not usage.

Time to first action, completion rate, independent completion.

A4. Isn't this a niche market?

Kerb ramps were built for a few hundred wheelchair users and are now used by everyone with a suitcase. Text messaging was built for deaf users. Captions were built for the hard of hearing and are how most people watch video in public. Accessibility features have a consistent history of outgrowing the population they were built for — and in the meantime, one in five students is not niche.

The curb-cut effect. Accommodations become markets.

A5. Are you diagnosing people?

No, and we are deliberate about it. Onboarding asks what tends to get in your way, not what you have been diagnosed with, and the result is a preference profile rather than a medical record. Nothing in the product asserts a condition about anyone.

A preference profile, not a diagnosis.

A6. Does this replace human disability support staff?

It extends their reach rather than replacing them. A support office typically has a handful of staff for thousands of students, so they triage. SETU handles the moment-to-moment friction of everyday digital tasks so that human time goes to the higher-level work only a person can do.

It extends the support office, it does not replace it.

B • BUSINESS MODEL AND BUYER

B1. Who actually pays?

Institutions, not students. A university licenses campus-wide at ₹199 per student per year; an employer buys per seat at ₹499 from an accessibility or DEI budget that already exists. Individuals get the local tools free forever, because those genuinely cost us nothing to serve.

Universities and employers, from budgets that already exist.

B2. Why campus-wide rather than per seat?

Because the buyer cannot identify the users. India registers 0.1% of its students as disabled against a far higher true prevalence, so a per-seat licence would only cover students already receiving support — the group that needs us least. Campus-wide is the only structure that reaches the rest, and it works commercially because most licensed students only ever use the tools that cost us nothing.

You cannot buy seats for people nobody has identified.

B3. How did you arrive at ₹199?

Bottom-up from cost. An active user costs ₹25 a month, activation in comparable campus tools runs around twelve per cent, so a five-thousand-student campus costs roughly ₹2.2 lakh a year to serve and prices at ₹9.95 lakh. That is a 78% gross margin, and the headroom is what funds the permanently free individual tier.

Cost-up, not market-down. 78% gross margin.

B4. Why is retail ₹299 a month when institutions pay ₹199 a year?

Because institutions buy five thousand seats of which about twelve per cent activate, so the effective price per active student is around ₹1,650 a year — roughly ₹138 a month. Against ₹299 retail that is a 2.2× volume discount, which is ordinary enterprise pricing.

Per active user it is a 2.2× discount, not 18×.

B5. Institutional sales cycles are long. How do you actually get in?

We do not start at central procurement. We start at the disability support office, which has a statutory mandate under the RPwD Act, a budget line, and more students than staff. A semester pilot is a far shorter decision than an IT purchase, and the pilot's outcome numbers are what unlock the second campus.

Disability support office first, procurement second.

B6. Realistically, what does year one look like?

Ten campuses and about one crore, with break-even at two. We would rather present a number we can hit in two semesters than a curve we cannot defend.

Ten campuses, ₹1 crore, break-even at two.

B7. What is the enterprise angle for a company like Capgemini?

Major employers have largely solved how to hire neurodivergent talent — JPMorgan's programme now spans more than forty roles across nine countries, and SAP, HPE and Microsoft have all reformed hiring to reach this group. What none of them has solved is making the internal tooling usable once those people arrive. That is a retention problem with a software answer, and it sits inside an existing DEI budget.

They solved hiring. Nobody solved the tooling.

C • COST, MARGIN AND SCALE

C1. What does this cost at a million users?

Most of a million users cost nothing, because the ten local tools never reach our servers. At twelve per cent monthly-active on AI features, a million licensed users is roughly ₹3.6 crore a year in inference and voice, against ₹19.9 crore of revenue at campus pricing.

The architecture is why the margin holds at scale.

C2. You are running on free model tiers. That is not a business.

Agreed, and we do not present it as one. We present it as evidence that the system degrades: the free tier is the top of a chain that ends in a deterministic offline engine. So we already know precisely what happens when it disappears — the ₹25 figure on our slide is that scenario, fully costed.

The free tier is a fallback chain, not a business model.

C3. Voice is your largest cost line. What if Sarvam raises prices?

Three answers. Clips are cached at both ends, so a repeated label is synthesised once. The browser's own voice is a working fallback at zero cost, and the engine switches to it automatically — even mid-passage. And text-to-speech is the most competitive segment in AI right now, so the price risk points down rather than up.

Cached, fallback-covered, and in a deflating market.

C4. What has this cost you to build so far?

Zero in infrastructure. Render, Vercel, MongoDB Atlas and OpenRouter free tiers, with a deterministic offline engine for when they rate-limit. Five students built four production surfaces for the cost of their time.

₹0 of infrastructure. Four surfaces. Five people.

C5. What if activation is much higher than twelve per cent?

Then margin compresses and the price moves — at twenty per cent activation, campus pricing goes to about ₹299 and the model still works. Activation is our least certain assumption and we would rather find out on one campus than on five hundred.

Our softest number, and the pilot is designed to test it.

C6. A 78% gross margin sounds too good. Justify it.

It is a normal software margin, and it is high specifically because the architecture keeps most users off our servers. The eighty-eight per cent of licensed students who only use the local reading tools cost us literally nothing. If we had built this cloud-first, the same product would carry a fraction of that margin.

The margin is a consequence of the architecture.

D • TECHNOLOGY AND ARCHITECTURE

D1. What is your stack, precisely?

Node and Express with MongoDB for the engine — fifty-three endpoints, nine cognitive modes. A Manifest V3 Chrome extension with fourteen content scripts and no build step. React with Vite for the web workspace. React Native with Expo for Android. Sarvam for speech in eleven Indian languages, and a model chain that starts on OpenRouter and falls through to Gemini and then to a deterministic offline engine.

Know this cold. Never inflate it.

D2. How do the local tools work with no server at all?

They are content scripts operating directly on the rendered DOM in the browser's own JavaScript engine — bio-reading weighting, focus mode, line focus, the reading ruler, auto-scroll, themes, and read-aloud through the browser's speech synthesis. No network call is involved in any of them, which is why they work offline and cost us nothing per use.

Content scripts on the rendered DOM. No network, no cost.

D3. Does it work on JavaScript-heavy single-page applications?

Yes, because we read the rendered DOM as the browser currently has it, not the source, so dynamically loaded content behaves exactly like static content. The extension runs on every http and https page with no per-site configuration.

We read the rendered DOM, so SPAs are not a special case.

D4. What is the latency?

Local tools are instant. The three-step path and the visual map render from the page in under a second and then upgrade in place when the model replies, typically two to four seconds later. Read-aloud is pipelined by sentence, so the first sentence starts in about a second instead of after the whole passage has synthesised.

Under a second locally, then it upgrades in place.

D5. What was your hardest technical problem?

Making AI features usable on rate-limited free models. A tool that shows a spinner and then an error is a tool nobody opens twice — so every AI feature now renders something real off the page in under a second and upgrades when the model replies. That constraint made the product better than an unlimited budget would have.

The constraint produced the render-then-upgrade pattern.

D6. How do you know it works? What is your test coverage?

A thirty-two point automated smoke test on the backend covering health probes, model round-trips, persistence, all nine modes, document OCR, streaming, and voice synthesis across all eleven languages — including an assertion that Hindi and Tamil responses keep their schema enums in English, because a model that helpfully translates a priority label silently breaks the UI. The extension has seven separate suites, none of which needs a bundler.

32 backend checks, 7 extension suites, and we test the failure modes.

D7. How do you handle eleven languages without the whole thing breaking?

One configuration file is the single source of truth for both halves — the speech language code and the instruction that makes the model answer in that language. Two details matter in practice: schema enum values stay in English so the client can compare them, and the sentence chunker knows the danda, because without it a Hindi paragraph reads as one enormous sentence and gets cut mid-clause.

One source of truth, English enums, and the chunker knows the danda.

E • AI SAFETY AND RELIABILITY

E1. Autonomous web agents hallucinate and break on real forms. How is yours different?

Three hard rules, all asserted in the test suite rather than promised in a slide: no step may target a control the page did not report, every submit step carries a confirmation requirement, and no personal data is ever invented to fill a field. The agent reads the page's semantic accessibility surface rather than guessing from pixels — the same approach Microsoft ships in Playwright MCP. It is an assisted execution engine, not an autonomous one.

Assisted execution, human-in-the-loop, asserted in tests.

E2. What if the page changes while the agent is mid-task?

The plan is validated against the controls the page currently reports, so a control that has vanished fails that step with a specific error code the UI keys its guidance off — it does not click something else instead. We would rather stop and tell the user than act on a stale plan.

It fails loudly rather than acting on stale state.

E3. How do you stop the AI from getting the content wrong?

We keep its job narrow and structural. It restructures content that is already on the page into a fixed JSON schema — it is never asked to generate new facts. Because everything renders from a schema, a malformed answer fails visibly instead of quietly producing something plausible and wrong.

Restructuring, not generating. Schema failures are visible.

E4. What about false positives — you interrupt someone who was working fine?

The detector is deliberately conservative, because a false positive interrupts someone who was fine and that is worse than a missed detection. Signals decay continuously, the detector only arms after a settling period, and once dismissed it will not re-trigger for several minutes. The prompt itself is a dismissible overlay, not a modal.

Conservative by design; a false positive is the worse error.

E5. Someone in crisis uses your reflective listener. What happens?

Crisis language is detected server-side before any model call, and the reply comes from a fixed, reviewed script carrying Tele-MANAS on 14416, KIRAN, AASRA and an international directory. It is never sampled from a model, never depends on the AI being reachable, and never scores points in the reward system. It is also deliberately not translated — a subtle translation error is unacceptable in that specific message, so it stays in English and leads with phone numbers, which need no translation.

Detected before the model, answered from a fixed script, never sampled.

F • PRIVACY, ETHICS AND COMPLIANCE

F1. You are reading behaviour and documents. What do you store, and where?

Identity is an anonymous per-device token — no name, no email, no account required. The overwhelm signals are computed in browser memory from cursor and scroll movement and are never transmitted; there is no camera, no biometrics and no keystroke logging. Journal entries in the reflective listener are written to a local-only key and are never mirrored to the database, and the API keys live server-side only so the browser never sees them.

Anonymous by default. The most sensitive data never leaves the device.

F2. How does that stand up to GDPR or an institutional DPA?

Our strongest position is data minimisation — we hold almost nothing personally identifying, which removes most of the surface an assessment would examine. For an institutional deployment we would expect to sign a data processing agreement, complete a VPAT, and add tenant-level deletion and export. Those are pilot deliverables and we would rather scope them with a real customer than guess at them now.

Minimisation first; DPA and VPAT are pilot deliverables.

F3. What accessibility standards do you meet?

WCAG 2.1 AA for structure and contrast, and Universal Design for Learning for presenting the same content in the form a given person processes best. The mobile client was checked against AA contrast, font scaling, screen-reader labels and forty-eight pixel touch targets. We have not yet had an independent audit, which is budgeted in year one.

WCAG 2.1 AA and UDL. Independent audit is budgeted, not done.

F4. Is the AI trained on user data?

No. We call third-party models through their APIs and route around any tier that reserves the right to train on payloads. We hold no training pipeline of our own, and given who our users are, we do not intend to build one.

No training on user data, and no ambition to.

F5. What is the risk of a tool like this being paternalistic?

It is a real risk and it drove several design decisions. Every intervention is offered and dismissible, never forced. The reward system is quiet by default and can be switched off entirely without losing progress. And the onboarding asks about preferences rather than diagnoses, because the moment a tool tells someone what they are, it has stopped being an accommodation.

Offered, never forced. Preferences, never diagnoses.

G • COMPETITION AND MOAT

G1. Microsoft Immersive Reader, Speechify, Texthelp all exist. What stops them crushing you?

They ship reading. Immersive Reader has been free and built into Edge and Office for years, and 95.9% of the web still fails accessibility standards — which tells you reading was never the binding constraint. Nobody has shipped the layer after reading: detecting when executive function stops, restructuring the task, acting on the page with confirmation, and carrying context across devices. And they will not build eleven Indian languages with a dyscalculia mode, because that market is too small for them and exactly right for us.

They solved reading. Reading was not the constraint.

G2. What genuinely stops a well-funded team copying this in six months?

Any one feature is copyable. What compounds is the combination — behavioural detection makes the intervention timely, a persistent cross-surface profile makes each one better than the last, structured task breakdown produces the outcome the institution is buying, and eleven Indian languages make it usable where we win first. The more durable moat is distribution: a support office that has run a semester pilot, has our numbers and has trained its staff does not re-procure casually.

The loop compounds, and the relationship is the real moat.

G3. Why would a university that already owns Read&Write buy you?

Because a university that owns a reading tool still has students who cannot start the form. We are not competing on text-to-speech. Several campuses will run both, and that is fine — we are a different line item solving a different failure.

Different category, different budget line, coexists happily.

G4. Is there anything defensible in the technology itself?

Honestly, the individually defensible piece is smaller than the product. What is hard is the accumulated judgement — the detector being conservative enough not to annoy people, the render-then-upgrade pattern that keeps AI features usable on rate-limited tiers, the schema enums staying in English so eleven languages do not break the UI. None of that is patentable and all of it took a clinician and several rebuilds to arrive at.

The moat is accumulated judgement plus distribution, not a patent.

H • TEAM, ROADMAP AND THE HARD ONES

H1. What is not built yet?

Three things, and we would rather say them than have you find them. There is no institutional dashboard, which a university will want by month two of a pilot. The app's own menus are still English-only — we change what SETU writes and says in eleven languages, not its own chrome. And we have no outcome data, only a working product.

No dashboard, English-only UI chrome, no outcome data yet.

H2. If you win today, what happens in ninety days?

A 250-student cohort with one university's disability support office, measuring the three metrics on our closing slide. The institutional dashboard, because it is what stands between a pilot and a renewal. And an internal trial with a Capgemini employee resource group if you will have us — a workplace cohort tests a different failure mode than a campus one.

250-student pilot, the dashboard, and an ERG trial.

H3. You all graduate this year. Why should we believe this continues?

It is a fair question and the honest answer is that a pilot changes the decision. Right now this is five final-year students with a running product and no obligation to anyone. One campus with a signed semester pilot, and a clinician already asking for a hospital build, is a materially different choice — which is exactly why our ask is a pilot rather than a prize.

The pilot is what converts this from a project into a company.

H4. What is the weakest part of this pitch?

That we are asking you to believe a product works before anyone has measured whether it does. Everything technical we have shown you is running and you can check it. The impact claim rests on an established intervention and one clinician's review — not on our own data, because we do not have any yet. That is the gap, and closing it is the only thing we are asking for.

Answer this one honestly. It is the highest-scoring answer in the set.

THE FIVE ANSWERS YOU MUST NOT IMPROVISE. Rehearse these five word-for-word until they are automatic. They are the ones where a fumble is expensive:

D1 (your stack) · E1 (agent safety) · F1 (what you store) · B2 (why campus-wide) · H4 (your weakest point).

WHEN YOU DO NOT KNOW. Say so, then convert it into a commitment. *"I don't have that number. I'd rather not guess — can I follow up with it after the session?"* That answer costs you almost nothing. A guess that turns out to be wrong costs you everything you said before it. Also agree in advance who answers what: one person on business and cost, one on architecture and safety, and neither interrupts the other. A team that visibly hands off is read as a team that has roles.

Pre-flight

PART 9

The last forty-eight hours, in order of what actually moves your score.

IN PRIORITY ORDER

	Action	Why it is at this priority
1	Email three disability NGOs asking for a letter of intent to pilot. Name one on slide 6.	Partnerships are a quarter of the published rubric and you currently score zero.
2	Correct the six number problems in Part 2.1 — especially the ₹199 chart and break-even.	A single contradiction found on stage contaminates every other figure.
3	Strike every unsupported claim from the master playbook and replace with Part 2.2.	"Show me the WASM module" is a normal question from this panel.
4	Rehearse the full seven minutes against a stopwatch, four times, standing up, out loud.	The most common way finalists lose is being cut off before the ask.
5	Record the demo video to the Part 7.4 screenplay. Caption it.	Insurance, and the hybrid demo needs it.
6	Rehearse the five must-not-improvise answers word for word.	Eight minutes of Q&A is longer than the pitch.
7	Run <code>npm run smoke</code> . Wake the engine. Test mirroring on the venue display.	Mechanical, but a dead demo erases a good pitch.

ON THE DAY — WHO DOES WHAT

Role	Responsibility
Speaker 1	Slides 1, 2, 3 and 6. Owns the story and the ask. In Q&A takes impact, business model, partnerships and the honest ones.
Speaker 2	Slide 4 (hands on keyboard) and slide 5. In Q&A takes architecture, cost, AI safety and privacy.
Third member	Sits with the laptop. Watches the clock, cues the fallback video, and does not speak unless directly asked. Silent competence reads well.
Everyone	When a teammate is answering, look at them, not at the jury and not at your notes. Never interrupt to add a detail — write it down and offer it only if there is time.

THE LAST THING.

You have a running product, a clinician who changed his mind about you, a cost structure that most funded startups would envy, and a closing line that reframes an entire category in nine words.

The only way you lose this is by claiming something you cannot show. Cut every one of those claims tonight, and then go and be exactly as good as you actually are.

Sources

PART 10

Every external figure used in this document and on your slides, with where it comes from. If a judge challenges a number, this is the page you want to have read.

Claim	Source
Tech4Positive Futures judged on partnership quality, scalability, feasibility and potential impact; global finale judged by Group Executive Committee members	Capgemini, <i>Tech4Positive Futures</i> — capgemini.com/in-en/about-us/csr/tech4positive-futures/
"Designed by and for humans"; useful, accessible and ethical	Aiman Ezzat, quoted on the same Capgemini page
8.4 million digital-inclusion beneficiaries since 2018; 955,000 in 2025; target of 10 million by 2030	Capgemini, <i>Digital Inclusion</i> — capgemini.com/about-us/csr/digital-inclusion/
95.9% of the top million home pages fail WCAG, up from 94.8%; 56.1 errors per page, up 10.1%; first reversal in seven years	WebAIM Million, 2026 — webaim.org/projects/million/
8% pooled prevalence of specific learning disorders in India (95% CI 4–11%), 6 studies, 8,133 children	Scaria, Bhaskaran & George, <i>Indian Journal of Psychological Medicine</i> , 2023
45.0 million students, 1,289 universities, 48,246 colleges, 54,080 students with benchmark disabilities	AISHE 2023–24, Ministry of Education, India
15–20% of the population is neurodivergent; 780 million people with dyslexia; 1 in 4 people with ADHD also have dyslexia	Your own research brief, drawing on the International Dyslexia Association, CDC and Cognassist
Assistive technology market \$21.4B (2025) → \$48.83B (2035), 8.6% CAGR	Vantage Market Research, <i>Assistive Technology Market</i>
Gemini 2.5 Flash at \$0.30 per million input tokens, \$2.50 per million output	Published model pricing
Sarvam Bulbul v2 at ₹15 per 10,000 characters; Saaras STT at ₹30 per hour	Sarvam AI pricing — docs.sarvam.ai
European Accessibility Act in force since 28 June 2025	Directive (EU) 2019/882
JPMorgan Chase Autism at Work spans 40+ roles across nine countries	JPMorgan Chase newsroom, <i>Proven Value: Autism at Work</i>
SAP, Hewlett Packard Enterprise and Microsoft reformed hiring to access neurodiverse talent, reporting gains in productivity, quality, innovation and engagement	Austin & Pisano, "Neurodiversity as a Competitive Advantage", <i>Harvard Business Review</i> , May 2017
The curb-cut effect and its examples	Level Access; Bureau of Internet Accessibility; widely documented in universal-design literature
53 endpoints · 9 cognitive modes · 11 languages · 14 content scripts · 32-point smoke test · Manifest V3 v3.2.0	Your repository: <code>backend/routes/apiRoutes.js</code> , <code>backend/config/languages.js</code> , <code>backend/scripts/smoke-test.js</code> , <code>chrome-extension/manifest.json</code>

Prepared 23 August 2026 for Team SETU · Capgemini Tech4Positive Futures Grand Finale & Nationals. Built against the final deck as submitted and against the repository at D:\PROJECTS\setu.